

Logistic Regression

Binary Classification

Binary Classification

- ▶ Consider a binary classification problem, i.e., $\mathcal{C} = \{\mathcal{C}_1, \mathcal{C}_2\}$.
- ▶ Given a particular vector \mathbf{x} , we wish to assign it to one of the two classes. To this end, we use Bayes' rule and calculate the relevant posterior probability:

$$\begin{aligned} P(\mathcal{C}_1|\mathbf{x}) &= \frac{P(\mathbf{x}|\mathcal{C}_1)P(\mathcal{C}_1)}{P(\mathbf{x})} \\ &= \frac{P(\mathbf{x}|\mathcal{C}_1)P(\mathcal{C}_1)}{P(\mathbf{x}|\mathcal{C}_1)P(\mathcal{C}_1) + P(\mathbf{x}|\mathcal{C}_2)P(\mathcal{C}_2)} \\ &= \frac{1}{1 + \frac{P(\mathbf{x}|\mathcal{C}_2) P(\mathcal{C}_2)}{P(\mathbf{x}|\mathcal{C}_1) P(\mathcal{C}_1)}} \\ &= \frac{1}{1 + \exp \left\{ -\log \left[\frac{P(\mathbf{x}|\mathcal{C}_1)}{P(\mathbf{x}|\mathcal{C}_2)} \right] - \log \left[\frac{P(\mathcal{C}_1)}{P(\mathcal{C}_2)} \right] \right\}} \end{aligned}$$

- It can be written in the form of the **logistic function**:

$$P(\mathcal{C}_1|\mathbf{x}) = \frac{1}{1 + e^{-\xi}},$$

where

$$\xi = \underbrace{\log \left[\frac{P(\mathbf{x}|\mathcal{C}_1)}{P(\mathbf{x}|\mathcal{C}_2)} \right]}_{\text{likelihood ratio}} + \underbrace{\log \left[\frac{P(\mathcal{C}_1)}{P(\mathcal{C}_2)} \right]}_{\text{prior ratio}}.$$

- Choose a particular form for the class-conditional density function $P(\mathbf{x}|\mathcal{C}_i)$.

Multivariate Gaussian Model

Assume that the class-conditional densities are multivariate Gaussian with identical covariance matrix $\mathbf{\Sigma}$:

$$P(\mathbf{x}|\mathcal{C}_i) = \frac{1}{(2\pi)^{\frac{D}{2}} |\mathbf{\Sigma}|^{\frac{1}{2}}} \exp \left\{ -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu}_i)^\top \mathbf{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}_i) \right\}.$$


After a simple calculation, we have

$$P(\mathcal{C}_1|\mathbf{x}) = \frac{1}{1 + e^{-(\mathbf{w}^\top \mathbf{x} + b)}},$$

where

$$\begin{aligned} \mathbf{w} &= \mathbf{\Sigma}^{-1} (\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2), \\ b &= \frac{1}{2} (\boldsymbol{\mu}_2 + \boldsymbol{\mu}_1)^\top \mathbf{\Sigma}^{-1} (\boldsymbol{\mu}_2 - \boldsymbol{\mu}_1) + \log \left[\frac{P(\mathcal{C}_1)}{P(\mathcal{C}_2)} \right]. \end{aligned}$$

Properties of Logistic Function



- ▶ $\sigma(\xi) \rightarrow 0$ as $\xi \rightarrow -\infty$.
- ▶ $\sigma(\xi) \rightarrow 1$ as $\xi \rightarrow \infty$.
- ▶ $\sigma(-\xi) = 1 - \sigma(\xi)$.
- ▶ $\frac{d}{d\xi} [\sigma(\xi)] = \sigma(\xi)\sigma(-\xi)$.

Maximum Likelihood Formulation

Logistic Regression

- ▶ Predict a binary output $y_n \in \{0, 1\}$ from an input \mathbf{x}_n .
- ▶ The logistic regression model has the form

$$y_n = \sigma(\mathbf{w}^\top \mathbf{x}_n) + \epsilon_n,$$

where

$$\sigma(\xi) = \frac{1}{1 + e^{-\xi}} = \frac{e^\xi}{1 + e^\xi}.$$

- ▶ Model input-output by a conditional Bernoulli distribution

$$P(y_n = 1 | \mathbf{x}_n) = \sigma(\mathbf{w}^\top \mathbf{x}_n).$$

- ▶ Discriminative model: Directly model $p(y|\mathbf{x})$.

Logistic Regression: MLE

- ▶ Given $\{(\mathbf{x}_n, y_n) \mid n = 1, \dots, N\}$, the likelihood is given by

$$\begin{aligned} p(\mathbf{y}|\mathbf{X}, \mathbf{w}) &= \prod_{n=1}^N p(y_n = 1|\mathbf{x}_n)^{y_n} (1 - p(y_n = 1|\mathbf{x}_n))^{1-y_n} \\ &= \prod_{n=1}^N \sigma(\mathbf{w}^\top \mathbf{x}_n)^{y_n} (1 - \sigma(\mathbf{w}^\top \mathbf{x}_n))^{1-y_n}. \end{aligned}$$

- ▶ Then log-likelihood function is given by

$$\mathcal{L} = \sum_{n=1}^N \log p(y_n|\mathbf{x}_n) = \sum_{n=1}^N \left\{ y_n \log \sigma_n + (1 - y_n) \log(1 - \sigma_n) \right\},$$

where $\sigma_n = \sigma(\mathbf{w}^\top \mathbf{x}_n)$.

- ▶ This is a nonlinear function of \mathbf{w} whose maximum cannot be computed in a closed form.
- ▶ **Iterative re-weighted least squares (IRLS)** is a popular algorithm, derived from Newton's method.

Newton's Method

Mathematical Preliminaries

- ▶ Gradient
- ▶ Hessian matrix
- ▶ Gradient descent/ascent
- ▶ Newton's method

Gradient

- ▶ Consider a real-valued function $f(\mathbf{x})$, that takes a real-valued vector $\mathbf{x} \in \mathbb{R}^D$ as an input,

$$f(\mathbf{x}) : \mathbb{R}^D \longrightarrow \mathbb{R}.$$

- ▶ The gradient of $f(\mathbf{x})$ is defined by

$$\begin{aligned}\nabla f(\mathbf{x}) &= \frac{\partial f(\mathbf{x})}{\partial \mathbf{x}} \\ &= \left[\frac{\partial f}{\partial x_1}, \dots, \frac{\partial f}{\partial x_D} \right]^\top.\end{aligned}$$

Hessian Matrix

If $f(\mathbf{x})$ belongs to the class \mathcal{C}^2 , the Hessian matrix \mathbf{H} is defined as the symmetric matrix with the (i, j) -element $\frac{\partial^2 f(\mathbf{x})}{\partial x_i \partial x_j}$,

$$\begin{aligned}\mathbf{H} &= \nabla^2 f(\mathbf{x}) \\ &= \left[\frac{\partial^2 f(\mathbf{x})}{\partial x_i \partial x_j} \right] \\ &= \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_D} \\ \vdots & & & \vdots \\ \frac{\partial^2 f}{\partial x_D \partial x_1} & \frac{\partial^2 f}{\partial x_D \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_D^2} \end{bmatrix}. \\ &= \frac{\partial}{\partial \mathbf{x}} \left[\frac{\partial f(\mathbf{x})}{\partial \mathbf{x}} \right]^\top \\ &= \frac{\partial}{\partial \mathbf{x}} [\nabla f(\mathbf{x})]^\top\end{aligned}$$

Gradient Descent/Ascent

- ▶ The gradient descent/ascent learning is a simple (first-order) iterative method for minimization/maximization.
- ▶ **Gradient descent**: iterative minimization

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} - \eta \left(\frac{\partial \mathcal{J}}{\partial \mathbf{w}} \right).$$

- ▶ **Gradient ascent**: iterative maximization

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} + \eta \left(\frac{\partial \mathcal{J}}{\partial \mathbf{w}} \right).$$

- ▶ **Learning rate**: $\eta > 0$

Newton's Method

The basic idea of Newton's method is to optimize the quadratic approximation of the objective function $\mathcal{J}(\mathbf{w})$ around the current point $\mathbf{w}^{(k)}$.

The second-order Taylor series expansion of $\mathcal{J}(\mathbf{w})$ at the current point $\mathbf{w}^{(k)}$ gives

$$\begin{aligned}\mathcal{J}_2(\mathbf{w}) = & \mathcal{J}(\mathbf{w}^{(k)}) + \left[\nabla \mathcal{J}(\mathbf{w}^{(k)}) \right]^\top (\mathbf{w} - \mathbf{w}^{(k)}) \\ & + \frac{1}{2} (\mathbf{w} - \mathbf{w}^{(k)})^\top \nabla^2 \mathcal{J}(\mathbf{w}^{(k)}) (\mathbf{w} - \mathbf{w}^{(k)}),\end{aligned}$$

where $\nabla^2 \mathcal{J}(\mathbf{w}^{(k)})$ is the Hessian of $\mathcal{J}(\mathbf{w})$ evaluated at $\mathbf{w} = \mathbf{w}^{(k)}$.

Differentiate this w.r.t. \mathbf{w} and set it equal 0, which leads to

$$\nabla \mathcal{J}(\mathbf{w}^{(k)}) + \nabla^2 \mathcal{J}(\mathbf{w}^{(k)}) \mathbf{w} - \nabla^2 \mathcal{J}(\mathbf{w}^{(k)}) \mathbf{w}^{(k)} = 0,$$

Thus we have


$$\mathbf{w} = \mathbf{w}^{(k)} - \left[\nabla^2 \mathcal{J}(\mathbf{w}^{(k)}) \right]^{-1} \nabla \mathcal{J}(\mathbf{w}^{(k)}).$$

Hence the Newton's method is of the form


$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} - \left[\nabla^2 \mathcal{J}(\mathbf{w}^{(k)}) \right]^{-1} \nabla \mathcal{J}(\mathbf{w}^{(k)}).$$

Remark: The Hessian $\nabla^2 \mathcal{J}(\mathbf{w}^{(k)})$ should be **positive definite** for all t .

Illustration of Newton's Method



(a) convex



(b) non-convex

[Figure source: Murphy's book]

Logistic Regression:

Algorithms

Logistic Regression: Gradient Ascent

The gradient ascent learning has the form

$$\mathbf{w}^{\text{new}} = \mathbf{w}^{\text{old}} + \eta \left(\frac{\partial \mathcal{L}}{\partial \mathbf{w}} \right).$$

Calculate the gradient

$$\frac{\partial \mathcal{L}}{\partial \mathbf{w}} = \sum_t \{y_n - \sigma(\mathbf{w}^\top \mathbf{x}_n)\} \mathbf{x}_n.$$

Hence, the gradient ascent update rule for \mathbf{w} is

$$\mathbf{w}^{\text{new}} = \mathbf{w}^{\text{old}} + \eta \sum_t \left\{ y_n - \sigma \left(\left(\mathbf{w}^{\text{old}} \right)^\top \mathbf{x}_n \right) \right\} \mathbf{x}_n.$$

Detailed Calculation of Gradient

Recall the log-likelihood

$$\mathcal{L} = \sum_{n=1}^N \left[y_n \log \sigma_n + (1 - y_n) \log(1 - \sigma_n) \right].$$

Calculate

$$\begin{aligned} \frac{\partial \mathcal{L}}{\partial \mathbf{w}} &= \sum_{n=1}^N \left[y_n \frac{\sigma'_n}{\sigma_n} \mathbf{x}_n + (1 - y_n) \frac{-\sigma'_n}{1 - \sigma_n} \mathbf{x}_n \right] \\ &= \sum_{n=1}^N \left[y_n \frac{\sigma_n(1 - \sigma_n)}{\sigma_n} - (1 - y_n) \frac{\sigma_n(1 - \sigma_n)}{1 - \sigma_n} \right] \mathbf{x}_n \\ &= \sum_{n=1}^N [y_n(1 - \sigma_n) - (1 - y_n)\sigma_n] \mathbf{x}_n \\ &= \sum_{n=1}^N (y_n - \sigma_n) \mathbf{x}_n. \end{aligned}$$

Detailed Calculation of Hessian

Calculate the Hessian:

$$\begin{aligned}\nabla^2 \mathcal{L} &= \frac{\partial}{\partial \mathbf{w}} [\nabla \mathcal{L}]^\top \\ &= \frac{\partial}{\partial \mathbf{w}} \left[\sum_{n=1}^N (y_n - \sigma_n) \mathbf{x}_n^\top \right] \\ &= \sum_{n=1}^N -\sigma_n(1 - \sigma_n) \mathbf{x}_n \mathbf{x}_n^\top.\end{aligned}$$

We set the objective function $\mathcal{J}(\mathbf{w})$ as the negative log-likelihood:

$$\mathcal{J}(\mathbf{w}) = -\mathcal{L}(\mathbf{w}) = -\sum_{n=1}^N \left[y_n \log \sigma_n + (1 - y_n) \log(1 - \sigma_n) \right].$$

Thus, the gradient and the Hessian are:

$$\begin{aligned} \nabla \mathcal{J}(\mathbf{w}) &= -\sum_{n=1}^N (y_n - \sigma_n) \mathbf{x}_n, \\ \nabla^2 \mathcal{J}(\mathbf{w}) &= \sum_{n=1}^N \sigma_n (1 - \sigma_n) \mathbf{x}_n \mathbf{x}_n^\top. \end{aligned}$$

Logistic Regression: IRLS

- ▶ Newton's update has the form

$$\Delta \mathbf{w} = - \underbrace{\left[\sum_{n=1}^N \sigma_n (1 - \sigma_n) \mathbf{x}_n \mathbf{x}_n^\top \right]^{-1}}_{\text{inverse of Hessian, } [\nabla^2 \mathcal{J}(\mathbf{w})]^{-1}} \underbrace{\left[- \sum_{n=1}^N (y_n - \sigma_n) \mathbf{x}_n \right]}_{\text{gradient, } \nabla \mathcal{J}(\mathbf{w})},$$

- ▶ Newton's update reduces to **iterative re-weighted least squares (IRLS)**:

$$\Delta \mathbf{w} = (\mathbf{X} \mathbf{S} \mathbf{X}^\top)^{-1} \mathbf{S} \mathbf{b},$$

where

$$\mathbf{S} = \begin{bmatrix} \sigma_1(1 - \sigma_1) & & 0 \\ & \ddots & \vdots \\ 0 & & \sigma_N(1 - \sigma_N) \end{bmatrix},$$
$$\mathbf{b} = \begin{bmatrix} \frac{y_1 - \sigma_1}{\sigma_1(1 - \sigma_1)} \\ \vdots \\ \frac{y_N - \sigma_N}{\sigma_N(1 - \sigma_N)} \end{bmatrix}.$$

Algorithm Outline

Algorithm 1 IRLS

Input: $\{(\mathbf{x}_n, y_n) \mid n = 1, \dots, N\}$

- 1: Initialize $\mathbf{w} = \mathbf{0}$ and $w_0 = \log(\bar{y}/(1 - \bar{y}))$
 - 2: **repeat**
 - 3: **for** $n = 1, \dots, N$ **do**
 - 4: Compute $\sigma_n = \sigma(\mathbf{w}^\top \mathbf{x}_n + w_0)$
 - 5: Compute $s_n = \sigma_n(1 - \sigma_n)$
 - 6: Compute $b_n = \frac{y_n - \sigma_n}{s_n}$
 - 7: **end for**
 - 8: Construct $\mathbf{S} = \text{diag}(s_{1:N})$
 - 9: Update $\mathbf{w} = (\mathbf{X}\mathbf{S}\mathbf{X}^\top)^{-1} \mathbf{S}\mathbf{b}$
 - 10: **until** convergence
 - 11: **return** \mathbf{w}
-

Generalized Linear Models

Generalized Linear Models

The GLM is a flexible generalization of ordinary linear regression that allows for response variables that have error distribution models other than a normal distribution.

The GLM consists of three elements:

1. A probability distribution from the exponential family.
2. A linear predictor $\eta = \mathbf{X}^\top \boldsymbol{\theta}$
3. A **link function** $g(\cdot)$ such that $\mathbb{E}[\mathbf{y}|\mathbf{X}] = \boldsymbol{\mu} = g^{-1}(\eta)$.

Logistic Regression: Generalized Linear Model

- ▶ Bernoulli distribution for a univariate binary random variable $y \in \{0, 1\}$ with mean μ ,

$$p(y|\mu) = \mu^y (1 - \mu)^{1-y}.$$

- ▶ **Logit transform (log-odds)**, $\theta = \mathbf{w}^\top \mathbf{x} = \log\left(\frac{\mu}{1-\mu}\right)$, leads to $[0, 1] \mapsto \mathbb{R}$.
- ▶ The **logistic function**, which is the **inverse-logit**, gives $\mu = \frac{1}{1+e^{-\theta}} = \sigma(\theta)$.
- ▶ Thus, we have

$$\begin{aligned} p(y|\theta) &= \sigma(\theta)^y (1 - \sigma(\theta))^{1-y} \\ &= \sigma(\theta)^y \sigma(-\theta)^{1-y} \end{aligned}$$

Multiclass Extension

Multiclass Logistic Regression (Multinomial Logistic Regression)

- ▶ Model input-output by a **softmax** transformation of **logits** $\theta_k = \mathbf{w}_k^\top \mathbf{x}_n$:

$$\begin{aligned} p(y_n = k | \mathbf{x}_n) &= \text{softmax}(\theta_k) \\ &= \frac{\exp\{\mathbf{w}_k^\top \mathbf{x}_n\}}{\sum_{j=1}^K \exp\{\mathbf{w}_j^\top \mathbf{x}_n\}} \end{aligned}$$

- ▶ Given $\mathbf{Y} \in \mathbb{R}^{K \times N}$ (each column $\mathbf{y}_n \in \mathbb{R}^K$ follows the 1-of- K coding) and $\mathbf{X} \in \mathbb{R}^{D \times N}$, the likelihood is given by

$$p(\mathbf{Y} | \mathbf{X}, \mathbf{w}_1, \dots, \mathbf{w}_K) = \prod_{n=1}^N \prod_{k=1}^K p(y_n = k | \mathbf{x}_n)^{Y_{k,n}}.$$

- ▶ The log-likelihood is given by

$$\mathcal{L} = \sum_{n=1}^N \sum_{k=1}^K Y_{k,n} \log [p(y_n = k | \mathbf{x}_n)].$$

- ▶ One can apply Newton's update to derive IRLS, as in logistic regression.